

Rodrigo Teixeira

Director of Data & AI · Data Platform Ownership | Production LLM & ML Systems | Applied Optimization

rodrigo@rodex.dev | +55 11 930212911 | São Paulo, Brazil (Remote) | linkedin.com/in/strodrigo |
github.com/strod | rodex.dev

PROFESSIONAL SUMMARY

Data and AI leader with **9+ years** across data engineering, data science and applied machine learning — and still hands-on. Own a company-wide data platform end to end: a nine-schema Redshift warehouse fed by 16 source systems, a 99-table Athena/Iceberg lakehouse, and the orchestration, governance and data-quality layers around both. Design and ship the production LLM systems built on top — agent-safety layers that let LLM agents operate on live systems, and natural-language interfaces to the warehouse itself. Set the platform's technical roadmap, including its migration off an end-of-life orchestrator, and lead engineers and analysts cross-functionally while partnering with executive, marketing and operations stakeholders on prioritisation and trade-offs. Engineering training applied commercially: Bayesian marketing mix modelling, convex-optimization budget allocation, and an advertising bid engine closing the loop from ingestion to automated bid writes.

PROFESSIONAL EXPERIENCE

Marketing Analytics Director

Filterbuy — Fort Lauderdale, FL (Remote)

Jan 2024 — Present

D2C and marketplace e-commerce in HVAC air filtration. Own the company-wide data platform, the advertising decision systems on it, and the AI tooling layer above both.

- Set the data platform roadmap and lead delivery cross-functionally with engineers and analysts, prioritising across marketing, finance, operations and manufacturing stakeholders and reporting outcomes to the executive team
- Own two warehouse platforms: a nine-schema Redshift warehouse (36-table star schema, isolated PII chains) fed by **16 source systems**, and a **99-table Athena/Iceberg lakehouse** ingested event-driven via S3→EventBridge→Lambda→MERGE
- Built the Amazon advertising decision platform end to end — 34 scheduled pipelines across 17 report types feeding a production bid engine (19,000+ lines of Python, 62 test suites) that writes bids through the Ads API behind a circuit breaker
- Designed the outer optimization loop: per-campaign Hill saturation curves fitted with bootstrap uncertainty, solved as a **convex program in cvxpy** under expected, risk-averse and chance-constrained modes to set profit-maximizing ACoS targets across ~300 campaigns per round
- Architected a **61-skill agent-safety library** across three ad platforms giving LLM agents one mutation contract — dry-run default, idempotency keys, risk gating, blast-radius estimation, approval and redacted audit (59 test suites)
- Led migration off an end-of-life ETL platform (**253 jobs**) onto serverless AWS orchestration; re-platformed a 13-hour job with recurring timeout failures onto Step Functions
- Built a **Bayesian marketing mix model** on Google Meridian over a purpose-built Iceberg lakehouse, trained on ECS Fargate, with six diagnostic guards and production alerting on fit quality and data freshness
- Established warehouse governance: a 36-table data dictionary with validated field definitions, a 16-source lineage register, and a 99-page agent-readable data-lake catalog
- Diagnosed a **43.8% marketing-attribution gap** (66,850 of 152,575 orders) and root-caused it to four checkout mechanisms, separating expected-by-design from fixable loss; drove the gap from ~53% to ~37%
- Built a predictive 24-month LTV model (RFM + XGBoost) wired to a server-side pixel for media optimization — **12% uplift in repeat purchase orders**
- Led the GA4 migration with client- and server-side GTM, uncovering **30%+ conversion undercounting** from ITP and ad blockers; implemented causal-impact frameworks using geo-holdouts and matched-market experiments
- **Cut reporting latency from 24 hours to under 1 hour** via real-time pipelines, and **reduced warehouse query costs 40%** through partitioning, clustering and schema redesign

Head of Data LATAM

IPG Mediabrands — São Paulo, Brazil

Jan 2022 — Dec 2023

- Led first-party data unification across Latin America, integrating CRM and offline sales into the Google advertising ecosystem; implemented server-side measurement across multiple LATAM markets
- Delivered Nestlé's "Speedboat" programme — technical integration of CRM, cloud and digital advertising platforms; partnered with enterprise stakeholders to turn infrastructure needs into investment cases

Data Manager / Tech Lead

MINT — São Paulo, Brazil

Jan 2020 — Dec 2021

- Built **AIRE**, a next-best-investment engine using Markowitz-style portfolio optimization to allocate media budget across channels and DSPs — **increased ROI by 25%+**
- Developed **Megalista**, privacy-safe ML-driven audience generation, plus attribution models measuring channel effectiveness and incremental value

Data Analytics Manager

WMcCann (McCann Worldgroup) — São Paulo, Brazil

Jan 2019 — Dec 2019

- Led a marketing mix modelling study for General Motors (**\$100M+ annual spend**) defining optimal annual budget allocation, with automated measurement pipelines in SQL and BigQuery

Data Scientist

Rede Globo — Rio de Janeiro, Brazil

Jan 2018 — Dec 2018

- Automated reporting for Gshow, Latin America's largest entertainment portal — **reduced manual reporting time by 80%**; built a churn optimization algorithm for content engagement

Data Science Consultant

Bridge Consulting — Rio de Janeiro, Brazil

Jan 2017 — Dec 2017

- Built target-stock calculation and expense outlier-detection algorithms for retail clients, with custom dashboards and automated reporting

TECHNICAL SKILLS

Data Engineering	Dimensional modelling & warehouse design, Redshift, BigQuery, Athena, Apache Iceberg, Glue, Matillion, Fivetran, Airflow, dbt, event-driven pipelines (Pub/Sub, EventBridge, SNS), Step Functions, data governance (dictionaries, lineage, contracts, quality assertions)
AI & LLM Engineering	Agent tool/skill layer design (dry-run, idempotency, risk gating, approval, audit, rollback), production LLM agents & tool-calling orchestration, Model Context Protocol (MCP), text-to-SQL over warehouses, prompt-injection hardening, multi-LLM routing (Anthropic, OpenAI, Bedrock, Vertex AI), RAG & vector search (FAISS)
Analytics & Experimentation	Marketing Mix Modelling (Google Meridian, Meta Robyn, Bayesian MCMC), convex & portfolio optimization (cvxpy, SciPy), Empirical Bayes smoothing, attribution (Markov chain, multi-touch), predictive LTV (XGBoost, SHAP), A/B testing, causal impact (geo-holdouts, matched markets), forecasting (Prophet, ARIMA), RFM segmentation, churn prediction
Cloud & Infrastructure	GCP (BigQuery, Cloud Run, Cloud Functions, Pub/Sub, Firestore, Vertex AI, Secret Manager), AWS (Redshift, Athena, Lambda, S3, Step Functions, EventBridge, DynamoDB, ECS Fargate, Bedrock), Terraform, Docker, Kubernetes, GitHub Actions CI/CD with Workload Identity Federation & OIDC, OIDC/SSO & JWT auth
BI & Visualization	Amazon QuickSight, Looker & LookML, Looker Studio, Tableau, Power BI, Streamlit, custom analytics front-ends (React, CloudFront/S3)
MarTech & Advertising	Amazon Ads (Reporting API, Marketing Stream, AMC, Ads API), Google Ads & GAQL, Meta, TikTok, Bing, Google Merchant Center, GA4, GTM (client & server-side), DV360, Campaign Manager 360, Vertex AI Search for Retail; Shopify, Klaviyo, Okendo, Recharge, Triple Whale
Languages	Python (pandas, PyArrow, SciPy, cvxpy, Prophet, boto3), SQL, TypeScript / Node.js, React, FastAPI, Zod, pytest / Vitest — Portuguese (native), English (fluent), Spanish (intermediate)

EDUCATION

MBA in Data Science & Analytics — USP (Universidade de São Paulo)

2020 – 2021 · Thesis: Machine Learning Applied to Audience Generation in Digital Marketing

M.Sc. in Industrial Engineering (*interrupted*) — COPPE-UFRJ

2018 – 2019 · Research: Non-Linear Programming Methods for Media Mix Modelling

B.Sc. in Mechanical Engineering — UFF (Universidade Federal Fluminense)

2010 – 2015

B.Sc. in Physics — UFF (Universidade Federal Fluminense)

2006 – 2010